

|   |
|---|
| Experiment No. 3  |
| Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model |
| Date of Performance:07/08/23  |
| Date of Submission:20/08/23   |

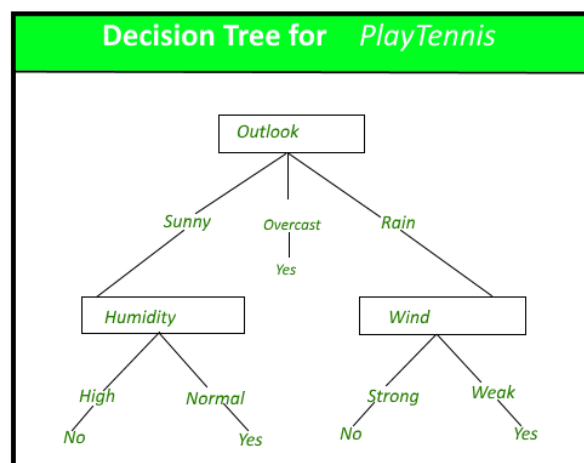


**Aim:** Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

**Objective:** To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

**Theory:**

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



**Dataset:**

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.



## Vidyavardhini's College of Engineering & Technology

### Department of Computer Engineering

---

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op- Inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic,



Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

**Code:**

**Conclusion:**

**1. Dealing with Categorical Attributes during Data Pre-processing:**

Managing Categorical Attributes During Data Pre-processing:

The values of the following categories columns are converted into unique number labels using label encoding. This numerical representation is required by many machine learning methods that need numerical input data. Encoding these categorical characteristics prepares the data for machine learning model training.

1. The term "workclass" refers to their job position.
2. The word "education" indicates their greatest degree of schooling.
3. "marital-status" indicates their marital status.
4. "occupation" displays their work functions.
5. "relationship" describes their familial situation.
6. "race" usually refers to their racial heritage.
7. The word "sex" denotes gender.
8. "native-country" frequently refers to their country of origin or citizenship. Certain columns are eliminated during data pre-processing in the code you gave.

**2. The following columns are specifically removed:**

1. Channel: This column is removed with the `data.drop(labels=(['Channel','Region']),axis=1,inplace=True)` function. The Channel column appears to have been deleted from the dataset.
2. Region: The Region column, like the Channel column, is discarded using the same line of code. This column is also deleted from the dataset. **Hyperparameter Tuning:**

The Decision Tree classifier is hyperparameter tuned in this code:



The Decision Tree classifier is built with a maximum depth of 5 specified: `DecisionTreeClassifier(max_depth=5)`. Because it influences the depth of the tree, this is a type of hyperparameter adjustment. This code sample, however, does not show a comprehensive hyperparameter tweaking procedure. In reality, more extensive approaches such as grid search or random search can be used to find the optimum hyperparameters. Only the `max_depth` is changed here.

### 3. Evaluation Metrics for Classification Models Confusion Matrix:

It accurately identified 4310 cases as negative (0) and 767 instances as positive (1), but it also projected 243 positives and 713 negatives incorrectly. Performance Metrics: The accuracy for positive predictions (1) is 0.76 lower than in Model 1, while the recall is 0.52 higher. This model has an F1-score of 0.62. This model has an overall accuracy of 0.84.

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.86      | 0.95   | 0.90     | 4553    |
| 1            | 0.76      | 0.52   | 0.62     | 1480    |
| accuracy     |           |        | 0.84     | 6033    |
| macro avg    | 0.81      | 0.73   | 0.76     | 6033    |
| weighted avg | 0.83      | 0.84   | 0.83     | 6033    |



**Vidyavardhini's College of Engineering & Technology**

Department of Computer Engineering


---

Importing lib

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Load the dataset

```
df = pd.read_csv('adult.csv')
df.head(10)
```



|   | age | workclass | fnlwgt | education    | education.num | marital.status | occupation        | rela |
|---|-----|-----------|--------|--------------|---------------|----------------|-------------------|------|
| 0 | 90  | ?         | 77053  | HS-grad      | 9             | Widowed        | ?                 | No   |
| 1 | 82  | Private   | 132870 | HS-grad      | 9             | Widowed        | Exec-managerial   | No   |
| 2 | 66  | ?         | 186061 | Some-college | 10            | Widowed        | ?                 | L    |
| 3 | 54  | Private   | 140359 | 7th-8th      | 4             | Divorced       | Machine-op-inspct | L    |
| 4 | 41  | Private   | 264663 | Some-college | 10            | Separated      | Prof-specialty    |      |
| 5 | 34  | Private   | 216864 | HS-grad      | 9             | Divorced       | Other-service     | L    |
| 6 | 38  | Private   | 150601 | 10th         | 6             | Separated      | Adm-              | L    |

Understanding Dataset

```
print ("Total Rows      : " ,df.shape[0])
dataset_row = df.shape[0]
print ("Total Columns   : " ,df.shape[1])
print ("\nFeatures      : \n" ,df.columns.tolist())
print ("\nMissing values : ", df.isnull().sum().values.sum())
print ("\nUnique values  : \n",df.nunique())

Total Rows      :  32561
Total Columns   :  15

Features :
['age', 'workclass', 'fnlwgt', 'education', 'education.num', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'capital.gain', 'capital.loss', 'hours.per.week', 'native.country', 'income']

Missing values :  0

Unique values :
age                73
workclass           9
fnlwgt            21648
education          16
education.num       16
marital.status      7
occupation          15
relationship        6
race                5
sex                 2
capital.gain        119
capital.loss         92
hours.per.week      94
native.country       42
income              2
dtype: int64
```

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass              32561 non-null  object
2   fnlwgt                 32561 non-null  int64
3   education              32561 non-null  object
4   education.num          32561 non-null  int64
```

```
5 marital.status 32561 non-null object
6 occupation     32561 non-null object
7 relationship   32561 non-null object
8 race           32561 non-null object
9 sex            32561 non-null object
10 capital.gain   32561 non-null int64
11 capital.loss   32561 non-null int64
12 hours.per.week 32561 non-null int64
13 native.country 32561 non-null object
14 income         32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
df.describe()
```

|       | age          | fnlwgt       | education.num | capital.gain | capital.loss | hours.per.week |
|-------|--------------|--------------|---------------|--------------|--------------|----------------|
| count | 32561.000000 | 3.256100e+04 | 32561.000000  | 32561.000000 | 32561.000000 | 32561.000000   |
| mean  | 38.581647    | 1.897784e+05 | 10.080679     | 1077.648844  | 87.303830    | 40.437456      |
| std   | 13.640433    | 1.055500e+05 | 2.572720      | 7385.292085  | 402.960219   | 12.347429      |
| min   | 17.000000    | 1.228500e+04 | 1.000000      | 0.000000     | 0.000000     | 1.000000       |
| 25%   | 28.000000    | 1.178270e+05 | 9.000000      | 0.000000     | 0.000000     | 40.000000      |
| 50%   | 37.000000    | 1.783560e+05 | 10.000000     | 0.000000     | 0.000000     | 40.000000      |
| 75%   | 48.000000    | 2.370510e+05 | 12.000000     | 0.000000     | 0.000000     | 45.000000      |
| max   | 90.000000    | 1.484705e+06 | 16.000000     | 99999.000000 | 4356.000000  | 99.000000      |

Missing Values

```
df_missing = (df=='?').sum()
print(df_missing)
```

```
age                0
workclass          1836
fnlwgt             0
education          0
education.num      0
marital.status     0
occupation         1843
relationship       0
race               0
sex                0
capital.gain       0
capital.loss       0
hours.per.week     0
native.country     583
income             0
dtype: int64
```

```
percent_missing = (df=='?').sum() * 100/len(df)
percent_missing
```

```
#dropping row having missing values from dataset
df = df[df['workclass'] != '?']
df = df[df['occupation'] != '?']
df = df[df['native.country'] != '?']
df.head()
```

|   | age | workclass | fnlwgt | education    | education.num | marital.status | occupation        | relationship  | race  | sex    | capital.gain | capital.loss |
|---|-----|-----------|--------|--------------|---------------|----------------|-------------------|---------------|-------|--------|--------------|--------------|
| 1 | 82  | Private   | 132870 | HS-grad      | 9             | Widowed        | Exec-managerial   | Not-in-family | White | Female | 0            | 0            |
| 3 | 54  | Private   | 140359 | 7th-8th      | 4             | Divorced       | Machine-op-inspct | Unmarried     | White | Female | 0            | 0            |
| 4 | 41  | Private   | 264663 | Some-college | 10            | Separated      | Prof-specialty    | Own-child     | White | Female | 0            | 0            |
| 5 | 34  | Private   | 216864 | HS-grad      | 9             | Divorced       | Other-service     | Unmarried     | White | Female | 0            | 0            |
| 6 | 38  | Private   | 150601 | 10th         | 6             | Separated      | Adm-clerical      | Unmarried     | White | Male   | 0            | 0            |

```
df_missing = (df=='?').sum()
print(df_missing)
```

```
age                0
workclass          0
fnlwgt             0
education          0
education.num      0
marital.status     0
```



```

occupation      0
relationship    0
race            0
sex             0
capital.gain    0
capital.loss    0
hours.per.week  0
native.country  0
income          0
dtype: int64

```

```

print("Total Rows after dropping rows : ",df.shape[0])
print("Numbers of rows drop: ", dataset_row -df.shape[0])

```

```

Total Rows after dropping rows : 30162
Numbers of rows drop: 2399

```

## Data Preparation

```
from sklearn import preprocessing
```

```

df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()

```

|   | workclass | education    | marital.status | occupation        | relationship  | race  | sex    | native.country |
|---|-----------|--------------|----------------|-------------------|---------------|-------|--------|----------------|
| 1 | Private   | HS-grad      | Widowed        | Exec-managerial   | Not-in-family | White | Female | U.S.           |
| 3 | Private   | 7th-8th      | Divorced       | Machine-op-inspct | Unmarried     | White | Female | U.S.           |
| 4 | Private   | Some-college | Separated      | Prof-specialty    | Own-child     | White | Female | U.S.           |

```

le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()

```

|   | workclass | education | marital.status | occupation | relationship | race | sex | native.country | income |
|---|-----------|-----------|----------------|------------|--------------|------|-----|----------------|--------|
| 1 | 2         | 11        | 6              | 3          | 1            | 4    | 0   | U.S.           | 38000  |
| 3 | 2         | 5         | 0              | 6          | 4            | 4    | 0   | U.S.           | 38000  |
| 4 | 2         | 15        | 5              | 9          | 3            | 4    | 0   | U.S.           | 38000  |
| 5 | 2         | 11        | 0              | 7          | 4            | 4    | 0   | U.S.           | 38000  |
| 6 | 2         | 0         | 5              | 0          | 4            | 4    | 1   | U.S.           | 38000  |

```

df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df['income'] = df['income'].astype('category')
df.head()

```

|   | age | fnlwgt | education.num | capital.gain | capital.loss | hours.per.week | workclass | education | marital.status | occupation | relationship  |
|---|-----|--------|---------------|--------------|--------------|----------------|-----------|-----------|----------------|------------|---------------|
| 1 | 82  | 132870 | 9             | 0            | 4356         | 18             | 2         | 11        | 6              | 3          | Not-in-family |
| 3 | 54  | 140359 | 4             | 0            | 3900         | 40             | 2         | 5         | 0              | 6          | Unmarried     |
| 4 | 41  | 264663 | 10            | 0            | 3900         | 40             | 2         | 15        | 5              | 9          | Own-child     |
| 5 | 34  | 216864 | 9             | 0            | 3770         | 45             | 2         | 11        | 0              | 7          | Married       |
| 6 | 38  | 150601 | 6             | 0            | 3770         | 40             | 2         | 0         | 5              | 0          | Married       |

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   age                 30162 non-null  int64
 1   fnlwgt              30162 non-null  int64
 2   education.num       30162 non-null  int64
 3   capital.gain        30162 non-null  int64
 4   capital.loss        30162 non-null  int64
 5   hours.per.week      30162 non-null  int64
 6   workclass           30162 non-null  int64

```

```

7  education      30162 non-null  int64
8  marital.status  30162 non-null  int64
9  occupation      30162 non-null  int64
10 relationship    30162 non-null  int64
11 race            30162 non-null  int64
12 sex             30162 non-null  int64
13 native.country  30162 non-null  int64
14 income          30162 non-null  category
dtypes: category(1), int64(14)
memory usage: 3.5 MB


```

```
sns.heatmap(df.corr(), cmap = 'RdGy')
```

```

<ipython-input-101-b22fcbbd6ef9>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future ve
sns.heatmap(df.corr(), cmap = 'RdGy')
<Axes: >

```



### Splitting dataset

```
from sklearn.model_selection import train_test_split
```

```

X = df.drop('income',axis=1)
X = X.drop('sex',axis=1)
y = df['income']

```

```
X.head()
```

|   | age | fnlwgt | education.num | capital.gain | capital.loss | hours.per.week | workclass | education | marital.status | occupation | relation |
|---|-----|--------|---------------|--------------|--------------|----------------|-----------|-----------|----------------|------------|----------|
| 1 | 82  | 132870 | 9             | 0            | 4356         | 18             | 2         | 11        | 6              | 3          |          |
| 3 | 54  | 140359 | 4             | 0            | 3900         | 40             | 2         | 5         | 0              | 6          |          |
| 4 | 41  | 264663 | 10            | 0            | 3900         | 40             | 2         | 15        | 5              | 9          |          |
| 5 | 34  | 216864 | 9             | 0            | 3770         | 45             | 2         | 11        | 0              | 7          |          |
| 6 | 38  | 150601 | 6             | 0            | 3770         | 40             | 2         | 0         | 5              | 0          |          |

```
y.head()
```

```

1    0
3    0
4    0
5    0
6    0
Name: income, dtype: category
Categories (2, int64): [0, 1]

```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.20)
```

## Appling Decision Tree Algo

```
from sklearn.tree import DecisionTreeClassifier
```

```
dt_default = DecisionTreeClassifier(max_depth=5)
dt_default.fit(X_train,y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5)
```

```
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
```

```
y_pred_default = dt_default.predict(X_test)
print("confusion matrix\n",confusion_matrix(y_test,y_pred_default))
print(classification_report(y_test,y_pred_default))
```

```
confusion matrix
[[4310  243]
 [ 713  767]]
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.86      | 0.95   | 0.90     | 4553    |
| 1            | 0.76      | 0.52   | 0.62     | 1480    |
| accuracy     |           |        | 0.84     | 6033    |
| macro avg    | 0.81      | 0.73   | 0.76     | 6033    |
| weighted avg | 0.83      | 0.84   | 0.83     | 6033    |

```
print("accuracy score: ",accuracy_score(y_test,y_pred_default))
```

```
accuracy score: 0.8415382065307475
```